

Machine Learning MIPT Introduction (seminar 1)

Ashuha Arseniy, Alexey Dral

ars.ashuha@gmail.com

February 9, 2016

Plan

1 Introduction

- Who we are?
- What Machine Learning is?

2 How to pass the course?

3 Optimization, Linear algebra and Probability and review

4 Data scientist tools (python and numpy review)

Introduction

1 Introduction

- Who we are?
- What Machine Learning is?

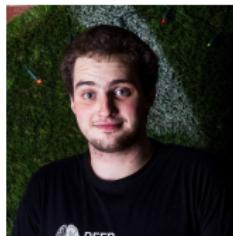
2 How to pass the course?

3 Optimization, Linear algebra and Probability and review

4 Data scientist tools (python and numpy review)

Course authors

- ▶ Arseniy Ashuha



Master student at MIPT, Data scientist at Rambler&CO

- ▶ Supervised by: Alexey Dral

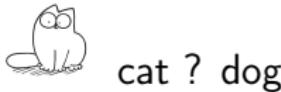


Assistant Prof. MIPT, Amazon

Machine learning in general

We have a visible variables X and we want to predict a hidden variables T.

- ▶ Classification



- ▶ Regression



- ▶ Phrase prediction



the guy sitting on chair and tune his guitar

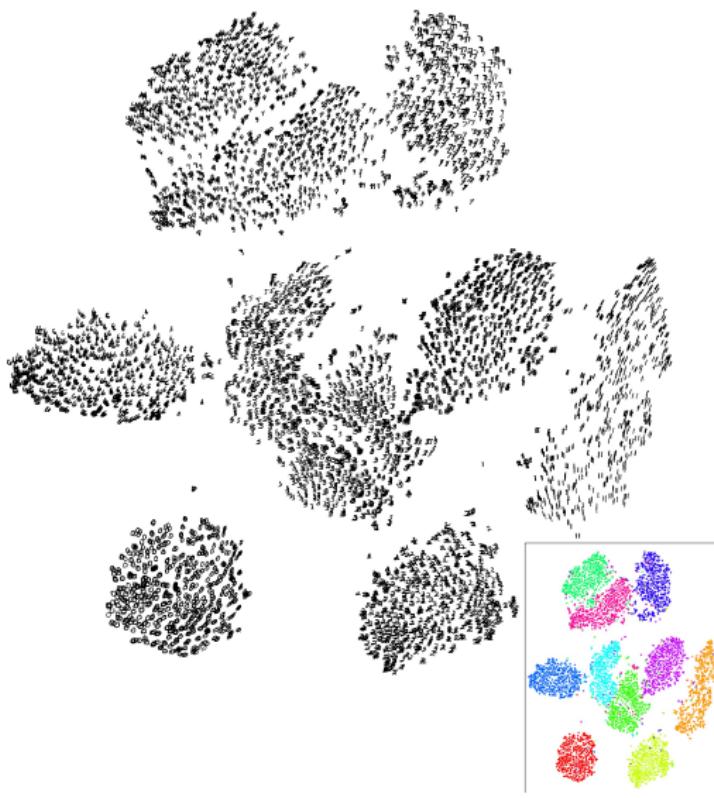
- ▶ Some art



) want to generate



Supervised / Unsupervised



How to pass the course?

1 Introduction

- Who we are?
- What Machine Learning is?

2 How to pass the course?

3 Optimization, Linear algebra and Probability and review

4 Data scientist tools (python and numpy review)

What will you learn after course

- ▶ How to look for ml problem in real live
- ▶ How to solve ml problem
 - ▶ What to do if train set is absent?
 - ▶ What to do if you have small/much data?
 - ▶ How to estimate quality of your model
- ▶ How to code it by yourself

How to pass the course?

- ▶ Practical HW (2 credits per HW) :
Bonus:
 - ▶ 1.0, 0.75, 0.5 for 1-3 positions
 - ▶ -0.2 each day after deadline, but not more -1.0
 1. KNN **picture classification code yourself**
 2. Trees, Forest **salary prediction code yourself**
 3. Linear, Dimension reduction **picture classification part code yourself**
 4. All Methods **age prediction**
- ▶ Practice class works (only bonus credits ☺):
 1. Bayesian classification
 2. Neural nets
- ▶ Theory credit (2 credit in time, 1 after)

$$\text{Result} = \text{sum} \cdot [\text{Theory credit} \geq 1.0]$$

course wiki: [MIPT ML 2016 Spring](#)

Linear algebra, probability and optimization review

1 Introduction

- Who we are?
- What Machine Learning is?

2 How to pass the course?

3 Optimization, Linear algebra and Probability and review

4 Data scientist tools (python and numpy review)

Optimization

- Usually we want to optimize function L – how lot error algorithm do

$$L(w, \text{data}) \rightarrow \min_w$$

- Optimize it by gradient descent

choose initial approximation of w and make small step by $-grad(w)$

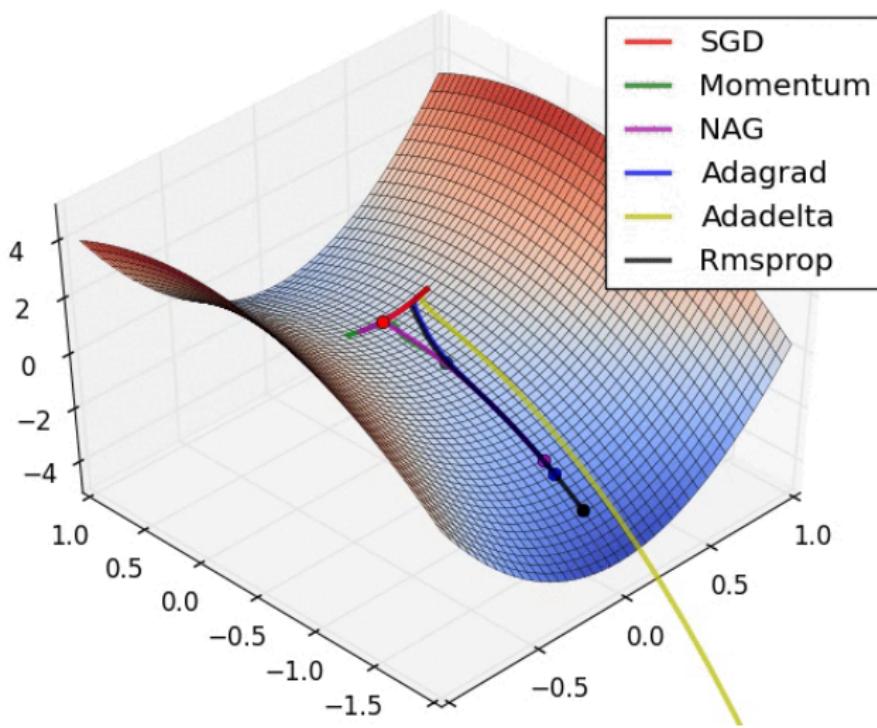
- Or stochastic gradient descent

$$grad(L) = \frac{d}{dw} L(w, \text{data}) \approx \frac{d}{dw} L(w, \text{random}(X))$$

- Linear restriction $r(w)$ e.g. $w_i \geq 0$, $\sum w_i = 1$ – w is equal probability for example

$$L(w, \text{data}) + \sum \lambda_i \cdot w_i + m \cdot (\sum w_i - 1) \rightarrow \min_w$$

Optimization



Linear algebra

- Sometimes we can express our algorithm in term of the Matrix (Tensor) multiplication – n^3 very fast operation 😊

"Dot Product"

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \times \begin{bmatrix} 7 & 8 \\ 9 & 10 \\ 11 & 12 \end{bmatrix} = \begin{bmatrix} 58 \end{bmatrix}$$

- T – train matrix, \hat{T} – predict matrix, we can use $\text{norms}(T - \hat{T})$

$$\|A\|_F = \sqrt{\sum_i \sum_j A_{ij}^2}$$

- If $\text{Det} = 0$ a lot of LA algorithms work poor

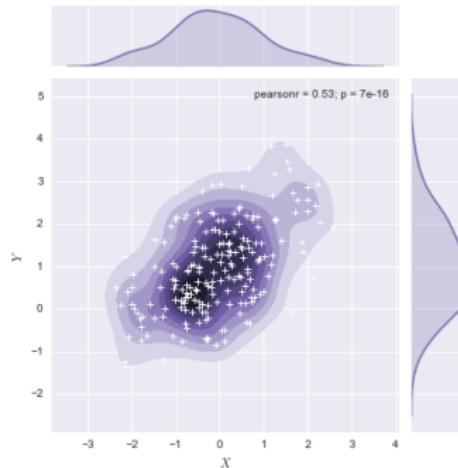
$$\det \begin{bmatrix} 1 & 2 \\ 1 & 2 + 1e^{-7} \end{bmatrix} \approx 0.000001 \quad \text{cond } \begin{bmatrix} 1 & 2 \\ 1 & 2 + 1e^{-7} \end{bmatrix} \approx 10000004$$

- Matrix calculus

$$L(W) = W_{11}^2 + W_{22} \quad \frac{d}{dW} L = \begin{bmatrix} 2 \cdot W_{11} & 0 \\ 0 & 1 \end{bmatrix}$$

Probability

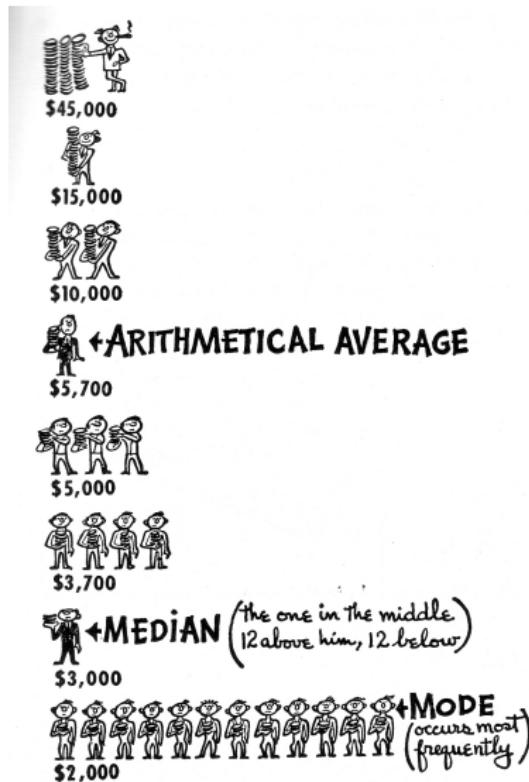
- ▶ Probability Mass Function + Conditional distribution + Correlation



- ▶ Bayesian rule, $p(T|X)$ – best approximator, but we cant find it general

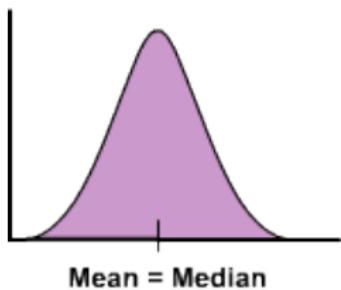
$$P(X|T) = \frac{P(T|X)P(X)}{P(T)}$$

Probability: Means

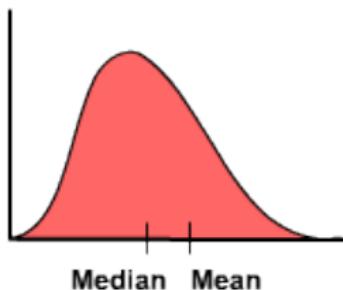


Probability: Means

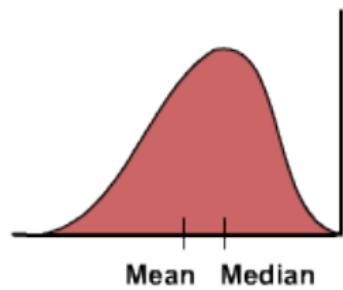
Symetric Distribution



Right-Skewed Distribution



Left-Skewed Distribution



Data scientist tools

1 Introduction

- Who we are?
- What Machine Learning is?

2 How to pass the course?

3 Optimization, Linear algebra and Probability and review

4 Data scientist tools (python and numpy review)

Python

1. Python – интерпретируемый язык, для быстрой разработки.
2. IPython

```
python2.7
MacBook-Air-arsl:~ arsl$ python2
Python 2.7.11 |Anaconda 2.4.0 (x86_64)| (default, Dec 6 2015, 18:57:58)
Type "copyright", "credits" or "license" for more information.

IPython 3.0.0 -- An enhanced Interactive Python.
?            -> Introduction and overview of IPython's features.
%quickref  -> Quick reference.
help       -> Python's own help system.
object?    -> Details about 'object', use 'object??' for extra details.

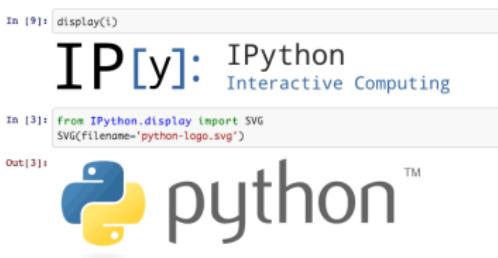
In [1]: print 123
123
```

3. Ipython Notebook + PyCharm + Anaconda

```
In [9]: display(i)
IP[y]: IPython
Interactive Computing

In [3]: from IPython.display import SVG
SVG(filename='python-logo.svg')

Out[3]:
```

A screenshot of an IPython Notebook interface. The top cell shows the command 'display(i)' and its output, which is the IPython logo. The bottom cell shows the command 'from IPython.display import SVG' and 'SVG(filename='python-logo.svg')', with its output being the Python logo itself.

Python: Useful packages

1. numpy/scipy – эффективная реализация массивов
2. pandas – работа с данными: чтение, join,
3. matplotlib/seaborn – мощная библиотека для графиков.



Настоятельные рекомендации по изучению Python stackoverflow

1. google-styleguide.googlecode.com/svn/trunk/pyguide.html
2. www.tutorialspoint.com/python/python_quick_guide.htm
3. O'REALLY Learning Python, 5th Edition

Solve ML Problem Stages

- ▶ Understand, what we want and how ml can help us – use your brain
- ▶ Get data – use your brain
- ▶ Read data – pandas
- ▶ Visualise and clean data – matplotlib, pandas, numpy
- ▶ Build model – ???
- ▶ Estimate quality – ???

Pandas

```
In [1]: import pandas as pd
```

```
In [2]: df = pd.read_csv('sal.csv')
```

```
In [3]: df
```

Out[3]:

	id	sex	salary	education
0	1	m	10000	high
1	2	m	5000	mid
2	3	f	7000	high
3	4	m	1000	none

```
In [4]: df.salary.mean()
```

Out[4]: 5750.0

```
In [12]: df.query('salary>=7000')
```

Out[12]:

	id	sex	salary	education
0	1	m	10000	high
2	3	m	7000	high

```
In [37]: df.sex = df.sex.apply(lambda x: 0 if x == 'm' else 1)
df.education = df.education.apply(
    lambda x: 1 if x == 'high' else 0)
```

```
In [41]: df
```

Out[41]:

	id	sex	salary	education
0	1	0	10000	1
1	2	0	5000	0
2	3	1	7000	1
3	4	0	1000	0

```
In [38]: X = df[['sex', 'education']]
T = df['salary']
```

```
In [39]: X.values
```

```
Out[39]: array([[0, 1],
                 [0, 0],
                 [1, 1],
                 [0, 0]])
```

```
In [40]: T.values
```

```
Out[40]: array([10000, 5000, 7000, 1000])
```

Links: [Pandas in 10 minutes](#), prof. Dyakonov materials: [pandas tutorial](#)

Python: Numpy is fast

```
In [1]: import numpy as np
```

```
In [2]: A, B = np.random.random((1000, 1000)), np.random.random((1000, 1000))
```

```
In [3]: %time C = A.dot(B)
```

```
CPU times: user 129 ms, sys: 5.8 ms, total: 135 ms
Wall time: 70.4 ms
```

```
In [4]: %time C = [[sum([A[i, j]*B[j, k] for j in xrange(1000)])
for i in xrange(1000)] for k in xrange(1000)]
```

```
CPU times: user 12min 28s, sys: 961 ms, total: 12min 29s
Wall time: 12min 29s
```

```
In [5]: a = np.random.randint(0, 10, 100000000)
```

```
In [6]: %time a = np.argsort(a)
```

```
CPU times: user 11 s, sys: 2.86 s, total: 13.9 s
Wall time: 14.7 s
```

Links: [Stanford numpy tutorial](#), [scipy tutorial](#)

matplotlib

```
In [1]: import numpy as np  
import matplotlib as pylab  
  
%pylab inline
```



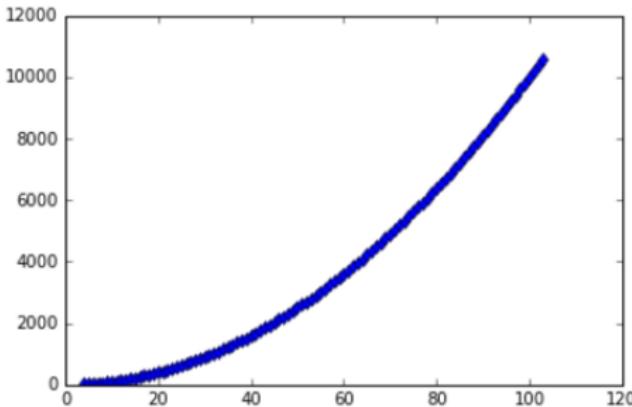
```
In [2]: x = np.arange(0, 100) + 3*np.random.normal()  
y = x**2 + 1
```



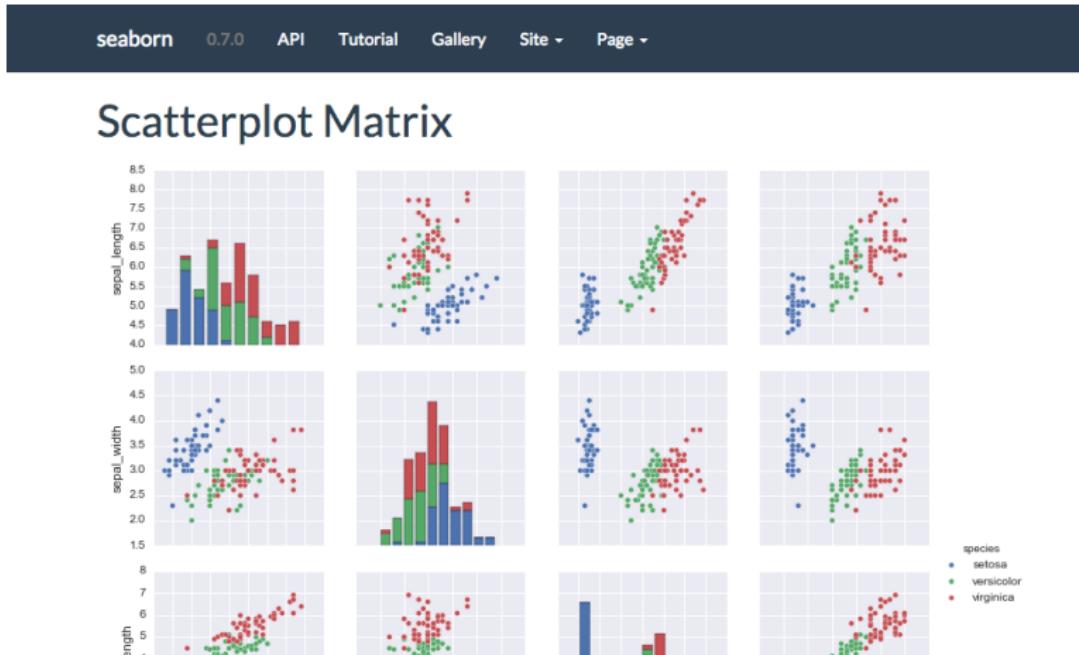
```
In [3]: pylab.plot(x, y, 'd')
```



```
Out[3]: [
```



seaborn



scipy

```
>>> from scipy.optimize import minimize, rosen, rosen_der
```

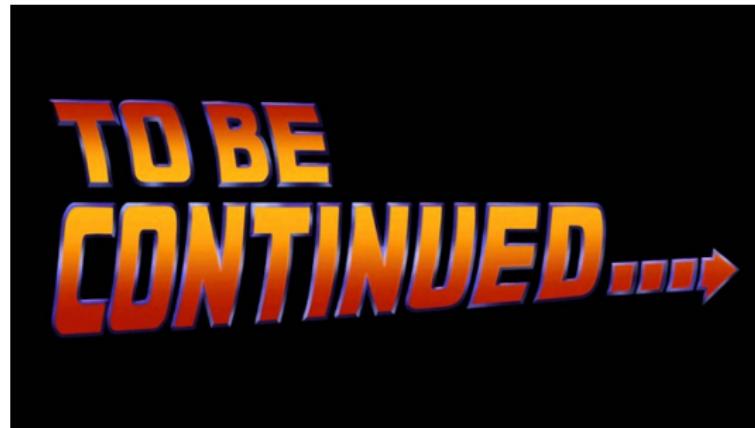
A simple application of the *Nelder-Mead* method is:

```
>>> x0 = [1.3, 0.7, 0.8, 1.9, 1.2]
>>> res = minimize(rosen, x0, method='Nelder-Mead')
>>> res.x
[ 1.  1.  1.  1.  1.]
```

Now using the *BFGS* algorithm, using the first derivative and a few options:

```
>>> res = minimize(rosen, x0, method='BFGS', jac=rosen_der,
...                  options={'gtol': 1e-6, 'disp': True})
Optimization terminated successfully.
      Current function value: 0.000000
      Iterations: 52
      Function evaluations: 64
      Gradient evaluations: 64
>>> res.x
array([ 1.  1.  1.  1.  1.])
```

learn model, estimate quality



Administrative

<https://goo.gl/UvtIO9>

- ▶ HW0: Install python, try libs, try wiki + piazza
- ▶ KNN HW will be announced next week